

# Sasidhar Mopuru

Data & AI Platform Engineer Databricks Certified

Bengaluru, India | sasidharmopuru@gmail.com | +91 63029 60712 | LinkedIn | GitHub | Portfolio

## PROFESSIONAL SUMMARY

Data & AI Platform Engineer with 2+ years building and validating production-grade streaming data pipelines and lakehouse systems at Accenture. Databricks Certified Data Engineer Associate and Microsoft DP-700. Core stack: PySpark, Kafka, Delta Lake, Databricks, Python, SQL. Delivered 90+ production PySpark ETL jobs, achieved 99.5% pipeline uptime, and reduced deployment time by 40%.

## CERTIFICATIONS

|                                                                                       |                                                                      |
|---------------------------------------------------------------------------------------|----------------------------------------------------------------------|
| Databricks Data Engineer Associate — Databricks, Dec 2024                             | DP-700: Data Engineering with Microsoft Fabric — Microsoft, Dec 2025 |
| Databricks PySpark Streaming Training — Accenture, Dec 2023 - Feb 2024                | Google Data Analytics Professional — Coursera, 2023                  |
| NPTEL Management Information System (MIS) — Elite · 73% · IIT Kharagpur, Jul–Oct 2022 |                                                                      |

## TECHNICAL SKILLS

|                   |                                                                                                                                 |
|-------------------|---------------------------------------------------------------------------------------------------------------------------------|
| Data Engineering  | PySpark, Apache Spark, Delta Lake, Apache Kafka, Azure Event Hubs, ETL/ELT, SQL, Data Modeling, Schema Validation, Data Quality |
| Cloud & Platforms | Databricks, Microsoft Fabric, Azure, AWS (S3, EMR Serverless, MSK Serverless)                                                   |
| DevOps & MLOps    | Docker, Terraform, GitLab CI/CD, GitHub Actions, Linux, pytest                                                                  |
| AI / LLM          | RAG, OpenAI, ChromaDB, Vector Databases, Sentence-Transformers, FastAPI, Streamlit                                              |
| Languages         | Python, SQL, YAML, JSON                                                                                                         |

## EXPERIENCE

Accenture — Data Engineer Feb 2024 – Present | Bengaluru, India

- Delivered 90+ production PySpark ETL jobs** across 4 supply-chain sub-domains, supporting 67 SDPs and 120+ CDPs through Python development, JSON configuration, and DDL development. (PySpark, SQL, Python, Git)
- Reduced deployment time by 40%** by developing configuration-driven pipeline definitions and reusable Python utilities used by the CDP/SDP platform and consumed by existing CI/CD workflows. (PySpark, CI/CD, JSON, Python)
- Achieved 99.5% pipeline uptime** through modular PySpark pipelines with error handling, retry logic, schema validation, and data quality checks. (PySpark, Kafka, Delta Lake)
- Maintained 95%+ overall test coverage** across 60+ pytest suites with mocked components and integration patterns. (pytest, unittest, Python)
- Improved data processing latency by 30%** through partitioning, caching, and performance tuning on production pipelines. (PySpark, Delta Lake, SQL)
- Owned end-to-end data quality and platform validation** for 3–4 sprint releases as the primary validation resource, validating schemas, tags, record counts, primary keys, business hash keys, and duplicate records across Kafka → Stage → Raw → HAST/CDP layers and preparing SQL-based reconciliation evidence for clean production sign-off. (SQL, Kafka, Data Quality)
- Developed Kafka consumer and validation workflows** for real-time event processing, DDL scripts, and validation/reconciliation queries for SDP/CDP layers, and monitored/troubleshoot production pipelines. (Kafka, Python, SQL, PySpark)

## PROJECTS

### Cloud-Native Streaming Data Platform

github.com/Sasireddy001/Cloud-data-platform

- Impact:** Exactly-once streaming with Delta Lake checkpointing, modular Terraform modules for dev/prod environments
- Highlights:** Multi-environment IaC, PySpark Structured Streaming, schema validation

### RAG Document QA Chatbot

github.com/Sasireddy001/rag-document-qa

- Impact:** Dense vector retrieval with response caching to reduce redundant LLM API calls
- Highlights:** PDF/DOCX/TXT ingestion, modular LLM backends, pytest tests with mocked components

### Kafka → PySpark → Delta Pipeline

github.com/Sasireddy001/Kafka-pyspark-delta-pipeline

- Impact:** 31k–45k rows/sec on a 4-core laptop, 95%+ pytest coverage, fault-tolerant exactly-once checkpointing
- Highlights:** JSON schema enforcement, automatic retry with exponential backoff, AWS Terraform deployment

## EDUCATION

B.Tech in Computer Science and Engineering — Graduated with Distinction

CGPA 7.99/10

JEE Mains — 94.14 percentile (Top 6% nationally among 1M+ candidates)