SASIDHAR MOPURU

Data Platform Engineer | Databricks Certified | PySpark | Kafka | Delta Lake | Python
Bengaluru, India | +91 63029 60712 | sasidharmopuru@gmail.com
LinkedIn | GitHub | Portfolio

SUMMARY

Data Platform Engineer with 2.7 years of experience building production-grade data pipelines, streaming platforms, and lakehouse systems at Accenture. Proven track record delivering 90+ PySpark ETL jobs across 67 source data products and 120+ core data products, with 99.5% uptime and 95%+ test coverage. Databricks DE Associate and DP-700 (Microsoft Fabric) certified. Expanding into cloud-native data platforms, RAG/LLM applications, and platform engineering.

EXPERIENCE

Data Engineer - Accenture | *Feb 2024 - Present* | *Bengaluru, India*

- **Delivered** 90+ production PySpark ETL jobs across 4 end-to-end supply-chain sub-domains, supporting **67 SDPs** and **120+ CDPs**.
- **Reduced deployment time by 70%** by building a configuration-driven Core Data Product (CDP) platform using JSON job definitions, Docker, and parent-child GitLab CI/CD pipelines.
- **Achieved 99.5% pipeline uptime** and reduced production incidents by 60% through modular Python utilities, schema validation, retry mechanisms, and comprehensive data quality checks.
- **Operated 95%+ overall test coverage** across 60+ pytest suites with mocked components and integration patterns.
- **Cut manual deployment effort by 60%** by containerizing Spark jobs and automating multi-environment (dev/sit/perf/prod) releases.
- **Reduced data processing latency by 30%** through partitioning, caching, and performance tuning on production pipelines.
- **Secured credential management** using HashiCorp Vault and enterprise certificate stores, meeting data governance and NDA standards.

PROJECTS

Cloud-Native Streaming Data Platform | GitHub
Azure Event Hubs | Databricks | ADLS Gen2 | Delta Lake | Terraform

- **Designed** a fully Terraform-managed streaming platform with multi-environment (dev/prod) modules and GitHub Actions CI/CD.
- **Processed** events through Azure Event Hubs into PySpark Structured Streaming with watermark-based deduplication and Delta Lake checkpointing.
- **Achieved 99.9% data accuracy** and cut infrastructure setup time by 50% through modular IaC.

Production-Style Kafka PySpark Delta Pipeline | GitHub
Apache Kafka | PySpark | Delta Lake | Databricks | GitHub Actions

- **Ingested** JSON events from Kafka, enforced schemas, and wrote exactly-once to Delta Lake using checkpointing and idempotent writes.
- **Benchmarked** 31k-45k rows/sec on a 4-core laptop for 100k-1M row workloads.
- **Maintained 90%+ pytest coverage** with an in-memory Spark fixture and continuous integration.

RAG Document QA Chatbot | GitHub
FastAPI | Streamlit | ChromaDB | OpenAI | Sentence-Transformers

- **Built** a retrieval-augmented generation application with document ingestion, chunking, dense vector retrieval, and LLM answer generation.
- **Achieved 85% retrieval accuracy** and reduced API calls by 70% through response caching.
- **Shipped** 85% pytest coverage with mocked LLM/embedding calls, FastAPI endpoints, and a Streamlit chat UI.

EDUCATION

B.Tech in Computer Science and Engineering - *CGPA 7.99*
JEE Mains - **94.14 Percentile (Top 6% nationally)**

CERTIFICATIONS

- DP-700: Implementing Data Engineering Solutions using Microsoft Fabric - Microsoft (2024)
- Databricks Certified Data Engineer Associate - Databricks (2024)
- Databricks PySpark Streaming Training (8 Weeks) - Accenture (2024)
- Google Data Analytics Professional Certificate - Coursera (2023)
- NPTEL MIS - IIT Madras (2023)

SKILLS

- **Languages:** Python, SQL
- **Data Engineering:** PySpark, Apache Spark, Delta Lake, Apache Kafka, ETL/ELT, Data Modeling, Data Quality, Schema Validation
- **Cloud & Platforms:** Databricks, Microsoft Fabric, Azure Event Hubs, ADLS Gen2, AWS (S3, EC2, Lambda), GCP (BigQuery, GKE)
- **DevOps:** GitLab CI/CD, GitHub Actions, Docker, Kubernetes, Terraform, HashiCorp Vault, Linux
- **AI / LLM:** RAG, LLM, FastAPI, Streamlit, ChromaDB, Vector Databases, OpenAI, Sentence-Transformers
- **Testing:** pytest, unittest, mocking, test-driven development