

Sasidhar Mopuru

Data & AI Platform Engineer Databricks Certified

Bengaluru, India | sasidharmopuru@gmail.com | +91 63029 60712 | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

PROFESSIONAL SUMMARY

Data & AI Platform Engineer with 2+ years building production-grade streaming data products and lakehouse systems at Accenture. Databricks Certified Data Engineer Associate and Microsoft DP-700. Core stack: PySpark, Kafka, Delta Lake, Databricks, Docker, Terraform. Delivered 5+ enterprise data products, achieved 99.5% pipeline uptime, and reduced deployment time by 40%.

CERTIFICATIONS

Databricks Data Engineer Associate — Databricks, Dec 2024	DP-700: Data Engineering with Microsoft Fabric — Microsoft, Dec 2025
Databricks PySpark Streaming Training — Accenture, Dec 2023 - Feb 2024	Google Data Analytics Professional — Coursera, 2023
NPTEL Management Information System (MIS) — Elite · 73% · IIT Kharagpur, Jul-Oct 2022	

TECHNICAL SKILLS

Data Engineering	PySpark, Apache Spark, Delta Lake, Apache Kafka, Azure Event Hubs, ETL/ELT, SQL, Data Modeling, Schema Validation, Data Quality
Cloud & Platforms	Databricks, Microsoft Fabric, Azure, AWS (S3, EC2, Lambda), GCP (BigQuery, GKE), SingleStore
DevOps & MLOps	Docker, Kubernetes, Terraform, GitLab CI/CD, GitHub Actions, HashiCorp Vault, Linux, pytest
AI / LLM	RAG, OpenAI, ChromaDB, Vector Databases, Sentence-Transformers, FastAPI, Streamlit
Languages	Python, SQL, YAML, JSON

EXPERIENCE

Accenture — Data Engineer Feb 2024 – Present | Bengaluru, India

- Reduced deployment time by 40% by building a configuration-driven CDP platform with PySpark, Docker, and GitLab CI/CD. (PySpark, Docker, CI/CD, Databricks)
- Delivered 5+ enterprise data products across batch and streaming architectures with end-to-end SDP and CDP data flows. (Azure, AWS, Delta Lake)
- Achieved 99.5% pipeline uptime through modular PySpark pipelines with error handling, retry mechanisms, and schema validation. (PySpark, Kafka, Delta Lake)
- Reduced data processing latency by 30% through partitioning, caching, and performance optimization strategies. (PySpark, Delta Lake, SQL)
- Implemented secure credential management with HashiCorp Vault and enterprise certificate stores; achieved 95%+ test coverage across pipeline components.
- Owned end-to-end data quality and platform validation for 3-4 sprint releases, validating schemas, tags, record counts, primary keys, business hash keys, and duplicate records across Kafka → Stage → Raw → HAST/CDP layers and preparing SQL-based reconciliation evidence for clean production sign-off. (SQL, Kafka, Data Quality, Platform Validation)

PROJECTS

Cloud-Native Streaming Data Platform

- github.com/Sasireddy001/Cloud-data-platform
- Impact: 99.9% data accuracy, exactly-once streaming, modular Terraform reducing setup time by 50%
 - Highlights: Multi-environment IaC, auto-scaling, spot-instance cost optimization, schema evolution

RAG Document QA Chatbot

- github.com/Sasireddy001/rag-document-qa
- Impact: 85% retrieval accuracy, 70% fewer API calls through intelligent caching
 - Highlights: PDF/DOCX/TXT ingestion, modular LLM backends, comprehensive pytest coverage

Kafka → PySpark → Delta Pipeline

- github.com/Sasireddy001/Kafka-pyspark-delta-pipeline
- Impact: 1M+ events/hour, 95%+ pytest coverage, fault-tolerant exactly-once checkpointing
 - Highlights: JSON schema enforcement, automatic retry with exponential backoff, AWS Terraform deployment

EDUCATION

B.Tech in Computer Science and Engineering — Graduated with Distinction CGPA 7.99/10
JEE Mains — 94.14 percentile (Top 6% nationally among 1M+ candidates)