

Sasidhar Mopuru

Data & AI Platform Engineer Databricks Certified

Bengaluru, India | sasidharmopuru@gmail.com | +91 63029 60712 | LinkedIn | GitHub | Portfolio

PROFESSIONAL SUMMARY

Data & AI Platform Engineer with 2+ years building production-grade streaming data products and lakehouse systems at Accenture. Databricks Certified Data Engineer Associate and Microsoft DP-700. Core stack: PySpark, Kafka, Delta Lake, Databricks, Docker, Terraform. Delivered 5+ enterprise data products, achieved 99.5% pipeline uptime, and reduced deployment time by 40%.

CERTIFICATIONS

Databricks Data Engineer Associate — Databricks, Dec 2024	DP-700: Data Engineering with Microsoft Fabric — Microsoft, Dec 2025
Databricks PySpark Streaming Training — Accenture, Dec 2023 - Feb 2024	Google Data Analytics Professional — Coursera, 2023
NPTEL Management Information System (MIS) — Elite · 73% · IIT Kharagpur, Jul–Oct 2022	

TECHNICAL SKILLS

Data Engineering	PySpark, Apache Spark, Delta Lake, Apache Kafka, Azure Event Hubs, ETL/ELT, SQL, Data Modeling, Schema Validation, Data Quality
Cloud & Platforms	Databricks, Microsoft Fabric, Azure, AWS (S3, EC2, Lambda), GCP (BigQuery, GKE), SingleStore
DevOps & MLOps	Docker, Kubernetes, Terraform, GitLab CI/CD, GitHub Actions, HashiCorp Vault, Linux, pytest
AI / LLM	RAG, OpenAI, ChromaDB, Vector Databases, Sentence-Transformers, FastAPI, Streamlit
Languages	Python, SQL, YAML, JSON

EXPERIENCE

Accenture — Data Engineer Feb 2024 – Present | Bengaluru, India

- Delivered 90+ production PySpark ETL jobs across 4 supply-chain sub-domains, supporting 67 SDPs and 120+ CDPs through Python development, JSON configuration, and DDL development. (PySpark, SQL, Python, Git)
- Reduced deployment time by 40% by developing configuration-driven pipeline definitions and reusable Python utilities used by the CDP/SDP platform and consumed by existing CI/CD workflows. (PySpark, CI/CD, JSON, Python)
- Achieved 99.5% pipeline uptime through modular PySpark pipelines with error handling, retry logic, schema validation, and data quality checks. (PySpark, Kafka, Delta Lake)
- Maintained 95%+ overall test coverage across 60+ pytest suites with mocked components and integration patterns. (pytest, unittest, Python)
- Improved data processing latency by 30% through partitioning, caching, and performance tuning on production pipelines. (PySpark, Delta Lake, SQL)
- Owned end-to-end data quality and platform validation for 3–4 sprint releases as the primary validation resource, validating schemas, tags, record counts, primary keys, business hash keys, and duplicate records across Kafka → Stage → Raw → HAST/CDP layers and preparing SQL-based reconciliation evidence for clean production sign-off. (SQL, Kafka, Data Quality)
- Developed Kafka consumer and validation workflows for real-time event processing, DDL scripts, and validation/reconciliation queries for SDP/CDP layers, and monitored/troubleshoot production pipelines. (Kafka, Python, SQL, PySpark)

PROJECTS

Cloud-Native Streaming Data Platform

github.com/Sasireddy001/Cloud-data-platform

- Impact: 99.9% data accuracy, exactly-once streaming, modular Terraform reducing setup time by 50%
- Highlights: Multi-environment IaC, auto-scaling, spot-instance cost optimization, schema evolution

RAG Document QA Chatbot

github.com/Sasireddy001/rag-document-qa

- Impact: 85% retrieval accuracy, 70% fewer API calls through intelligent caching
- Highlights: PDF/DOCX/TXT ingestion, modular LLM backends, comprehensive pytest coverage

Kafka → PySpark → Delta Pipeline

github.com/Sasireddy001/Kafka-pyspark-delta-pipeline

- Impact: 1M+ events/hour, 95%+ pytest coverage, fault-tolerant exactly-once checkpointing
- Highlights: JSON schema enforcement, automatic retry with exponential backoff, AWS Terraform deployment

EDUCATION

B.Tech in Computer Science and Engineering — Graduated with Distinction

CGPA 7.99/10

JEE Mains — 94.14 percentile (Top 6% nationally among 1M+ candidates)